

3

## References

- ❑ Deep Learning, Ian Goodfellow, Yoshua Bengio, Aaron Courville
- ❑ Neural Networks and Learning Machines, Simon Haykin
- ❑ Pattern Recognition and Machine Learning, Christopher M. Bishop
- ❑ Deep Learning with Python - François Chollet
- ❑ Hands-On Machine Learning with Scikit-Learn and TensorFlow
- ❑ TensorFlow Deep Learning Cookbook
- ❑ Reinforcement Learning with TensorFlow: A Beginner's Guide to Designing Self-learning Systems with TensorFlow and OpenAI Gym Sayon Dutta
- ❑ Hands-On Reinforcement Learning with Python: Master Reinforcement and Deep Reinforcement Learning Using OpenAI Gym and TensorFlow Sudharsan Ravichandiran
- ❑ Deep Reinforcement Learning Hands-On: Apply Modern RL Methods, with Deep Q-networks, Value Iteration, Policy Gradients, TRPO, AlphaGo Zero and More Maxim Lapan

11/18/2023

pra-sâmi

4

“ Theory exam- 40% weightage  
 Lab exam - 40% weightage  
 Internal exam - 20% weightage ”

### Evaluation method

11/18/2023

pra-sâmi

5

## Logistics



11/18/2023

pra-sâmi

6

## Logistics

- ❑ We encourage 'discussion groups':
  - ❖ Study groups
  - ❖ Whatsapp groups
- ❑ Expect you to complete your assignment **individually!**
- ❑ No group assignments unless stated otherwise
- ❑ Code is small part of it
- ❑ Pay special attention to **inline comments**
  - ❖ Comments should focus on what you were trying to implement




11/18/2023

pra-sâmi

7

## Journey So far



11/18/2023

pra-sâmi

8

## Uncertainty Everywhere



Goal:  
Delivering a passenger  
to the airport on time



- ❑ The agent forms a plan, lets say... A90,
  - ❖ Leave home 90 minutes before the flight departs
  - ❖ Driving at a reasonable speed
- ❑ Are you certain "Plan A90 will get us to the airport in time."?
  - ❖ Not in absolute sense but with some riders

11/18/2023

pra-sâmi

9

## Uncertainty Everywhere



Goal:  
Delivering a passenger  
to the airport on time



- ❑ How about other plans, such as A180,
  - ❖ Might increase the agent's belief that it will get to the airport on time,
  - ❖ But also increase the likelihood of a long wait
- ❑ Probability is an agent's measure of belief in some proposition — subjective probability.
- ❑ An agent's belief depends on its prior belief and what it observes.

11/18/2023

pra-sâmi

10

## Agent In Uncertain Environment

- ❑ Agents don't have complete knowledge about the world.
- ❑ Agents need to make (informed) decisions given their uncertainty.
- ❑ It isn't enough to assume what the world is like.
  - ❖ Example: wearing a seat belt.
- ❑ An agent needs to reason about its uncertainty.
- ❑ When an agent takes an action under uncertainty, it is gambling  $\Rightarrow$  probability



11/18/2023

pra-sâmi

11

## Overview

- ❑ Nature is a continuum where as math is discrete values
  - ❖ Old film based images were continuous painting of colors where as digital images are pixels
- ❑ Brain works differently than our mathematical computations
- ❑ Brain is highly complex, nonlinear and parallel computer
- ❑ Neural networks are supposed to be inspired from
- ❑ Highly generalized form, a Neural Network is a mathematical model that simulates the manner in which brain performs a task



All models are wrong... some models are useful!

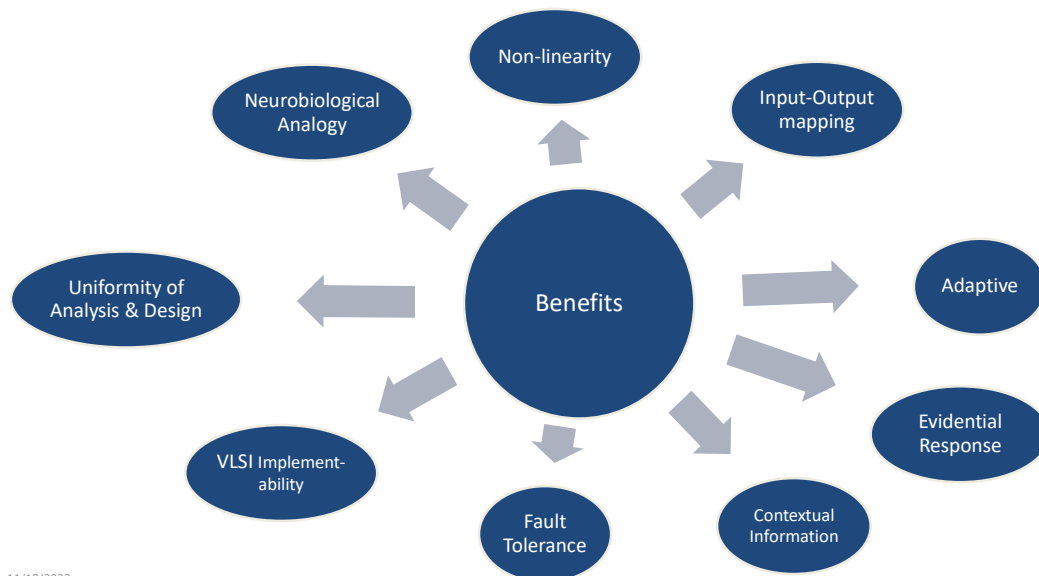
Is this how our brain works? Really!!

11/18/2023

pra-sâmi

12

## Benefits of Neural Networks



11/18/2023

pra-sâmi


13

## What's Being Played...



Can a Rule based system answer based on the data provided?

We as humans can easily guess!




11/18/2023

Time	Day	Type
9:00 AM	Weekday	News
11:00 AM	Weekday	K-Soaps
3:00 PM	Weekday	Soaps
5:00 PM	Weekday	Soaps
6:00 PM	Weekday	Cartoons
9:00 PM	Weekday	Sports, Movies
4:00 PM	Weekday	???

pra-sâmi


14

## AI vs ML vs Deep Learning



11/18/2023

- ❑ Used interchangeably
- ❑ AI is a broader concept, it includes basic AI to Deep learning.
- ❑ Machine learning: enabling Machines to Learn from the past incidents (available data).
- ❑ Deep Learning: One can say that it tries to copy information processing patterns found in the human brain



pra-sâmi

15

## Artificial Intelligence vs. Machine Learning

### Artificial Intelligence

- ❑ Create intelligent machines that can simulate Human thinking capability and behavior
- ❑ A man-made thinking power
- ❑ No pre-programming needed
- ❑ Algorithms which can work with their own "intelligence"
- ❑ Algorithms such as Reinforcement learning algorithm and deep learning neural networks being used in multiple places such as Siri, Google's AlphaGo, AI in Chess playing, etc.
- ❑ Based on capabilities, AI can be classified into three types:
  - ❖ Weak AI
  - ❖ General AI
  - ❖ Strong AI
- ❑ Currently, we are working with weak AI and general AI. The future of AI is Strong AI for which it is said that it will be more intelligent than humans (???)

### Machine Learning

- ❑ An application or subset of AI
- ❑ Allows machines to learn from data without being programmed explicitly
- ❑ Uses a massive amount of structured and semi-structured data
- ❑ It can work only on data it has seen
- ❑ For unknown cases it becomes unresponsive or unreliable
- ❑ Being used for online recommender system, for Google search algorithms, Email spam filter, Facebook Auto friend tagging suggestion, etc.
- ❑ It can be divided into three types:
  - ❖ Supervised learning
  - ❖ Unsupervised learning
  - ❖ Reinforcement learning

11/18/2023

pra-sâmi

16

## Deep Learning

- ❑ Large Neural Networks

"Using brain simulations, hope to:  
Make learning algorithms much better and easier to use,  
Make revolutionary advances in machine learning and AI,  
I believe this is our best shot at progress towards real AI."

- Andrew Ng

- ❑ Learning successive layers of increasingly meaningful representations
- ❑ Modern network contain hundreds of successive layers
- ❑ Successive layers are learned via "neurons" connected via neural network

Some concepts were inspired by how our brain works  
It is NOT a replica of human brain!!!

11/18/2023

pra-sâmi

17

## Deep Learning

- ❑ Why Deep Learning is more practical today?
  - ❖ Availability of large computing power
  - ❖ Availability of large datasets
- ❑ Most flavors of the old generations of learning algorithms, performance will plateau
- ❑ Deep learning that is scalable
  - ❖ Performance just keeps getting better as more and more data is fed
- ❑ Most value today is coming from supervised learning
- ❑ Eventually, we will see benefits of unsupervised learning

11/18/2023

pra-sâmi

18

## Deep Learning

- ❑ Usually a neural network contains
  - ❖ Input Layer
  - ❖ Hidden layers [ 1 ... n]
  - ❖ Output layer
- ❑ We may call network with 1 to 2 hidden layer as shallow
- ❑ Network with 10 or more layers as deep
  - ❖ No set demarcation!
- ❑ I guess, scientists just got excited when someone labeled them as deep network
- ❑ Intelligent software to automate routine tasks, understand speech or images, make diagnosis in medicine and support basic scientific research

11/18/2023


pra-sâmi




19

## How deep learning works...

- As the images are processed through the layers



- The representations are increasingly filtered, purified and distilled to make them more meaningful



11/18/2023

pra-sâmi

20

## What has been achieved so far

- Learn to see and hear... so natural to humans but elusive to machines earlier
- Image classification
- Speech recognition
- Handwriting recognition
- Writing style recognition (who was the author)
- Improved machine translation
- Text-to-speech conversion
- Digital assistants such as Google Now and Amazon Alexa
- Little autonomous driving
- Improved ad targeting, as used by Google, Baidu, and Bing
- Ability to answer natural-language questions
- Superhuman games playing: chess, go...

Still long way to go...  
Human-level general intelligence too far away...

Results on the web

11/18/2023

pra-sâmi

21

Neurons

11/18/2023

pra-sâmi

22

To play or not to play!

id	Dry Weather	Low Temp	Homework Done	Team Members	Equipment	Ground	Played
1	1	1	1	1	0	1	1
2	1	1	1	1	1	1	1
3	1	1	1	1	1	1	1
4	0	1	0	1	1	1	0
5	0	0	1	1	1	0	0
6	0	0	0	0	0	1	0

11/18/2023

pra-sâmi

23

## Neurons

### ❑ Features:

- ❖ Is it raining?
- ❖ Is it too hot?
- ❖ Have I completed my homework?
- ❖ Are sufficient players ready?
- ❖ Is cricket equipment ready?
- ❖ Is ground available?

### ❑ Depending on the feature values, you may get to play or not

### ❑ Features like homework and availability of ground can be considered as 'inhibitory'.

id	Dry Weather	Low Temp	Homework Done	Team Members	Equipment	Ground	Played
1	1	1	1	1	0	1	1
2	1	1	1	1	1	1	1
3	1	1	1	1	1	1	1
4	0	1	0	1	1	1	0
5	0	0	1	1	1	0	0
6	0	0	0	0	0	1	0

### ❑ Notes :

- ❖ Aggregator function is sum and threshold can be 3.
- ❖ Assign 0 or 1 if a parameter is in favor or not

Given sufficient data point, we can train an algorithm to make such simple decisions for us.

11/18/2023

pra-sâmi

24

## MP Neuron

- ❑ In 1943 Warren S. McCulloch, a neuroscientist, and Walter Pitts, a logician, published "A logical calculus of the ideas immanent in nervous activity" in the Bulletin of Mathematical Biophysics
- ❑ In this paper McCulloch and Pitts tried to understand how the brain could produce highly complex patterns by using many basic cells that are connected together
- ❑ These basic brain cells are called neurons, and McCulloch and Pitts gave a highly simplified model of a neuron in their paper
- ❑ The McCulloch and Pitts model of a neuron, which we will call an MCP neuron for short, has made an important contribution to the development of artificial neural networks -- which model key features of biological neurons


11/18/2023

pra-sâmi

25

## MP Neuron

- Neurons receive signals and produce a response
- In this model:
  - ❖ All inputs are binary i.e. [0,1]
  - ❖ Inputs are “inhibitory” or “excitatory”.
  - ❖ Inhibitory have maximum influence on the model
  - ❖ It has an aggregator ‘g’ and a function ‘f’
  - ❖ There is a threshold
  - ❖ If g is more than threshold,  $\hat{y} = 1$  else 0



□  $\hat{y} = 0$  if any  $x_i$  is inhibitory, else  $g(x) = \sum x_i$


□  $\hat{y} = 1$  if  $g(x) \geq \text{threshold}$  else  $\hat{y} = 0$

11/18/2023

pra-sâmi

26

## MP Neuron




id	Dry Weather	Low Temp	Homework Done	Team Members	Equipment	Ground	Sum	Played
1	1	1	1	1	0	1	5	1
2	1	1	1	1	1	1	6	1
3	1	1	1	1	1	1	6	1
4	0	1	0	1	1	1	4	0
5	0	0	1	1	1	0	3	0
6	0	0	0	0	0	1	1	0

11/18/2023

pra-sâmi

27

## MP Neuron



id	Dry Weather	Low Temp	Homework Done	Team Members	Equipment	Ground	Sum	Played
1	1	1	1	1	0	1	5	1
2	1	1	1	1	1	1	6	1
3	1	1	1	1	1	1	6	1
4	0	1	0	1	1	1	4	0
5	0	0	1	1	1	0	3	0
6	0	0	0	0	0	1	1	0

The logic is straight forward. Let's implement this model on a dataset.


11/18/2023

pra-sâmi

28

## Code Example1 – MP Neurons

- ❑ Need a dataset with plenty of features and binary output
- ❑ Breast Cancer dataset from scikit-learn
  - ❖ `data = sklearn.datasets.load_breast_cancer()`
- ❑ Its features are a continuous and we need binary
  - ❖ Use pandas `pd.cut` to bin the columns
  - ❖ `X_bin = X . apply ( pd.cut, bins=2, labels=[1,0])`
  - ❖ For `b` in range `[0, num_features+1]`
    - Sum it by row and compare with `b`
- ❑ Converted file is in the shared folder



11/18/2023

pra-sâmi

29



11/18/2023

*pra-sâmi*


ADDITIONAL MATERIAL

*pra-sâmi*

31


## Biological Neural Nets

- ❑ Pigeons as art experts (Watanabe et al. 1995)
- ❑ Experiment:
  - ❖ Pigeon in Skinner box
  - ❖ Present paintings of two different artists (e.g. Chagall / Van Gogh)
  - ❖ Reward for pecking when presented a particular artist (e.g. Van Gogh)



11/18/2023

pra-sâmi







34

## Biological Neural Nets

- ❑ Pigeons were able to discriminate between Van Gogh and Chagall
  - ❖ With 95% accuracy on train set (when presented with pictures they had been trained on)
  - ❖ Discrimination, still 85% successful for previously unseen paintings of the artists
- ❑ Pigeons do not simply memorise the pictures
- ❑ They can extract and recognise patterns (the 'style')
- ❑ They generalise from the already seen to make predictions
- ❑ This is what neural networks (biological and artificial) are good at (unlike conventional computer)



35

## Brain and Machine

### □ The Brain

- ❖ Pattern Recognition
- ❖ Association
- ❖ Complexity
- ❖ Noise Tolerance



### □ The Machine


- ❖ Calculation
- ❖ Precision
- ❖ Logic

11/18/2023

pra-sâmi


36

## The contrast in architecture



- The Von Neumann architecture uses a single processing unit;
  - ❖ Tens of millions of operations per second
  - ❖ Absolute arithmetic precision

- The brain uses many slow unreliable processors acting in parallel



11/18/2023

pra-sâmi

37

## The biological inspiration

- ❑ Features of the Brain
  - ❖ Ten billion ( $10^{10}$ ) neurons
  - ❖ On average, several thousand connections
  - ❖ Hundreds of operations per second
  - ❖ Die off frequently (never replaced)
  - ❖ Compensates for problems by massive parallelism
- ❑ The brain has been extensively studied by scientists
- ❑ Vast complexity prevents all but rudimentary understanding
- ❑ Even the behavior of an individual neuron is extremely complex
- ❑ Single “percepts” distributed among many neurons
- ❑ Localized parts of the brain are responsible for certain well-defined functions (e.g. vision, motion).



11/18/2023

pra-sâmi